



DESIGN, AUTOMATION & TEST IN EUROPE

14 – 15 March 2022 · on-site event  
16 – 23 March 2022 · online event

The European Event for Electronic  
System Design & Test

# FRL-FI: Transient Fault Analysis for Federated Reinforcement Learning-Based Navigation Systems

Zishen Wan<sup>1</sup>, Aqeel Anwar<sup>1</sup>, Abdulrahman Mahmoud<sup>2</sup>, Tianyu Jia<sup>3</sup>, Yu-Shun Hsiao<sup>2</sup>,  
Vijay Janapa Reddi<sup>2</sup>, Arijit Raychowdhury<sup>1</sup>

<sup>1</sup> Georgia Institute of Technology, <sup>2</sup> Harvard University, <sup>3</sup> Carnegie Mellon University

# Safety of Autonomous Navigation



- Swarm intelligence
- Specialized hardware accelerator
- Hardware fault
  - Low operational voltage
  - Technology scaling
- Traditional protection method
  - Hardware module redundancy

# Safety of Autonomous Navigation

- Swarm intelligence

How is resilience of swarm navigation system to hardware faults?

How do we detect and mitigate hardware faults?

Hardware protection

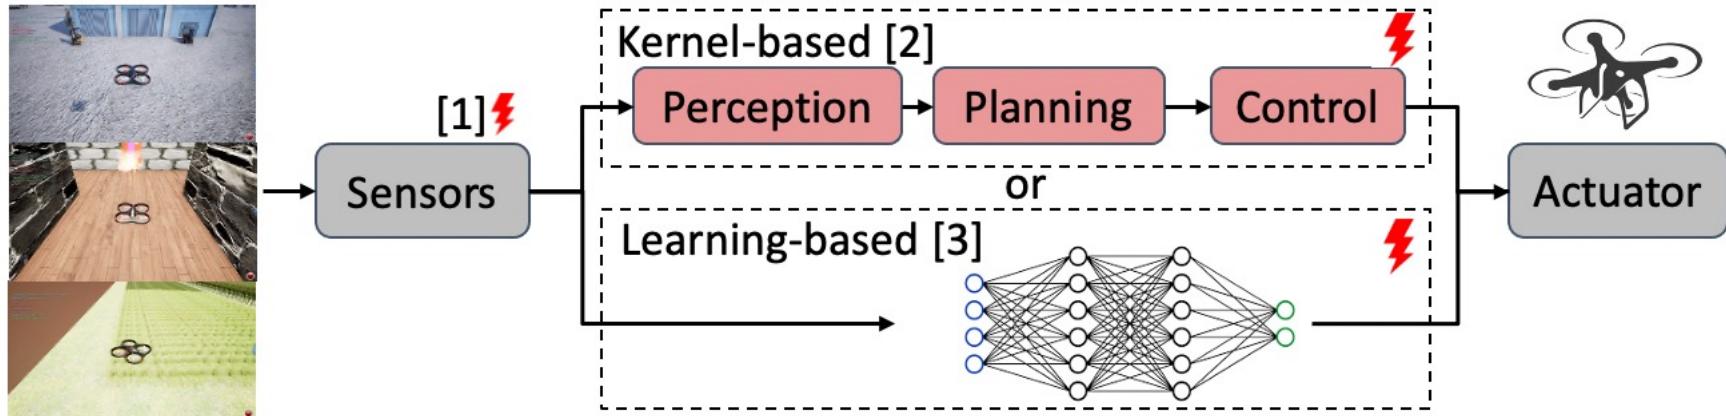
method

- Hardware module redundancy



# Related Work

- Reliability of single-agent autonomous system



[1] A. Toschi et al., NPC'19 [2] Y. Hsiao\*, Z. Wan\* et al., arXiv'21 [3] Z. Wan et al., DAC'21

How about reliability of multi-agent autonomous system (swarm intelligence)?

# Related Work

- **Fault characterization**
  - Neural network in supervised learning: PytorchFI, Ares, TensorFI

How about reliability of reinforcement learning-based (long-term decision making) system?

- **Fault mitigation**
  - Hardware redundancy-based method: DMR, TMR

Can we propose an application-aware lightweight protection method?

# This Work

## Transient Fault Analysis for Federated Reinforcement Learning (FRL) -Based Navigation Systems



Transient fault injection for FRL-based navigation systems



Transient fault characterization for FRL-based navigation systems



Transient fault mitigation for FRL-based navigation systems

# This Work

## Transient Fault Analysis for Federated Reinforcement Learning (FRL) -Based Navigation Systems



Transient fault injection for FRL-based navigation systems

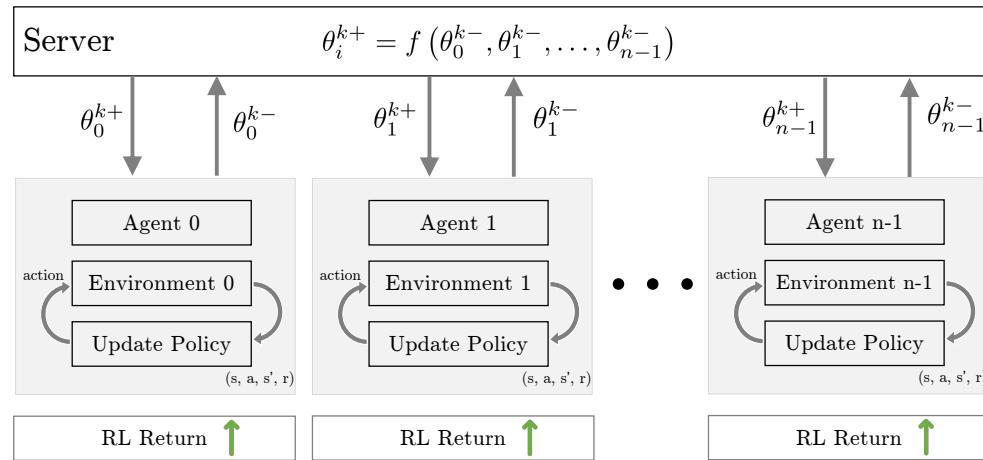


Transient fault characterization for FRL-based navigation systems



Transient fault mitigation for FRL-based navigation systems

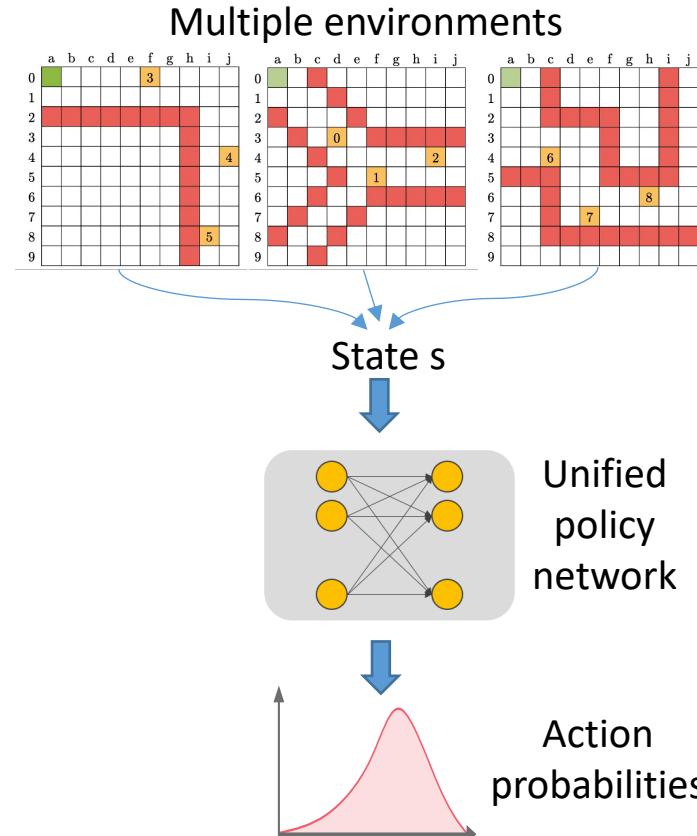
# Federated Reinforcement Learning (FRL) System



Policy trained on meta environments (simulator)  
--> transfer to real scenarios with fine-tuning

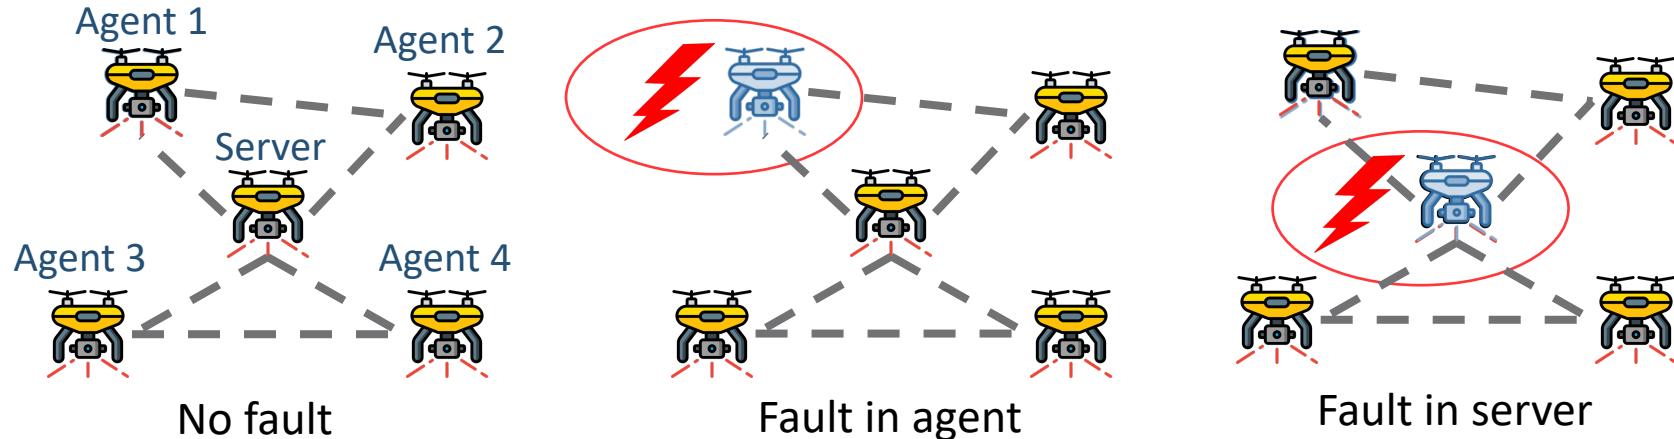
Training: phase with changing exploration-exploitation ratio

Inference: a greedy exploitation phase



# FRL with Faults

- **Fault Type: Transient fault**
  - Model: random bit-flip
- **Fault localization: memory**



- **Fault injection: static injection and dynamic injection**

# This Work

## Transient Fault Analysis for Federated Reinforcement Learning (FRL) -Based Navigation Systems



Transient fault injection for FRL-based navigation systems



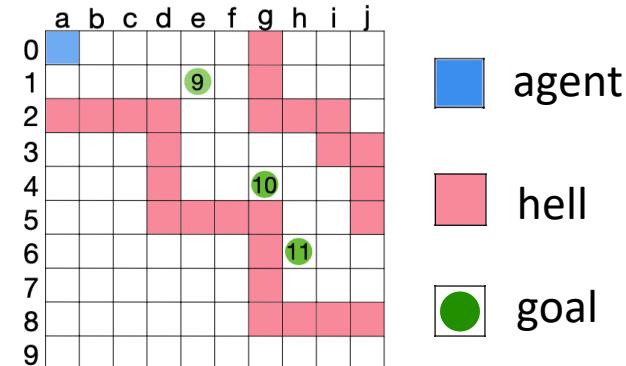
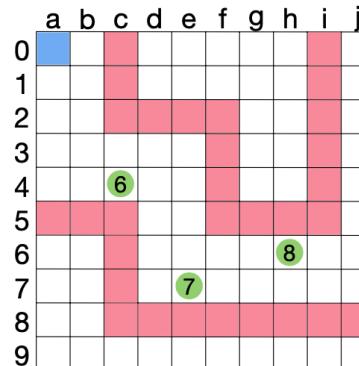
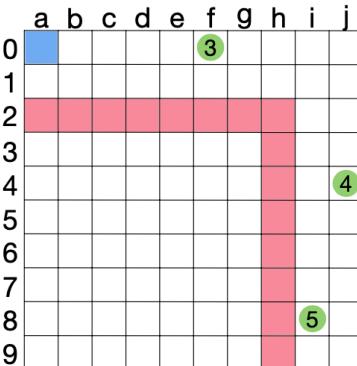
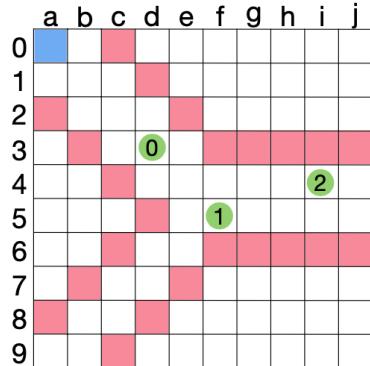
Transient fault characterization for FRL-based navigation systems



Transient fault mitigation for FRL-based navigation systems

# Grid-Based Navigation Problem (GridWorld)

- Goal: Start from the source position (█), reach the goal state (█) avoiding getting into hell (█)
- 12 Environments



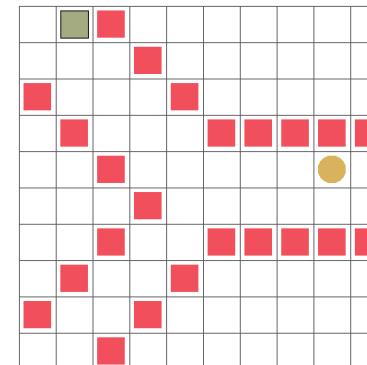
- Evaluation metric: average agents' success rate

# Grid-Based Navigation Problem (GridWorld)

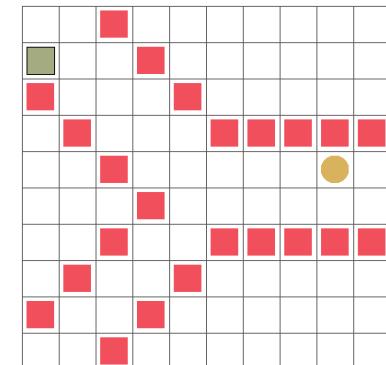
## Experiment Overview

- **Training (on-device fine-tuning)**
  - Transient fault impact
  - Server and agent comparison
  - Single and multi-agent comparison
  - Policy convergence
- **Inference**
  - Transient fault impact
  - Single and multi-agent comparison

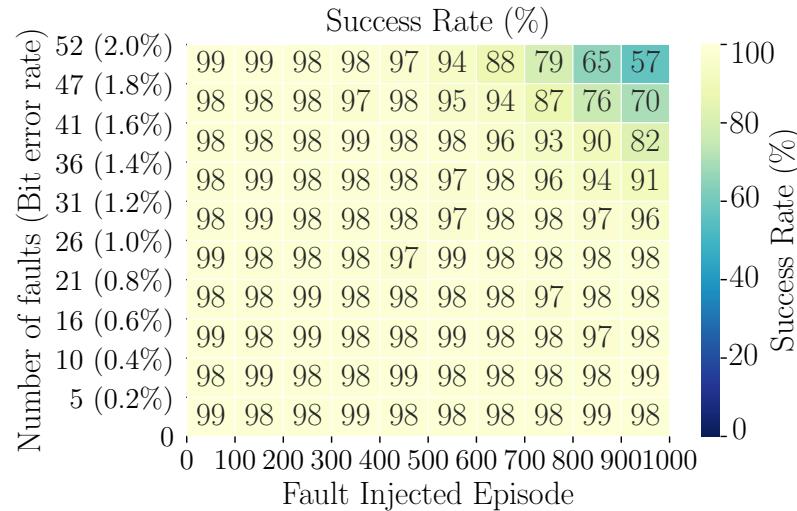
Policy performance  
in the free of fault



Policy performance  
in the presence of fault

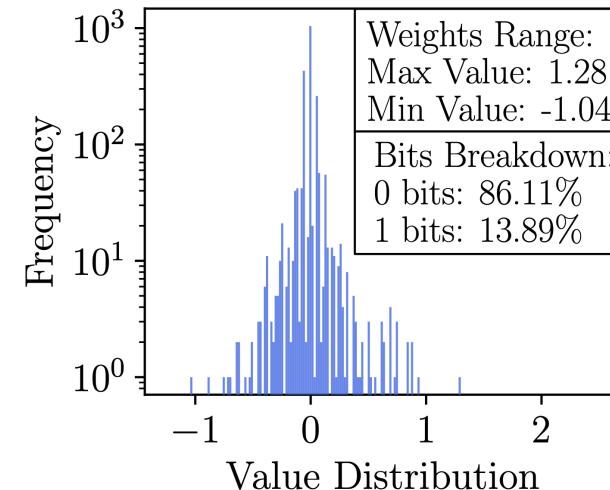
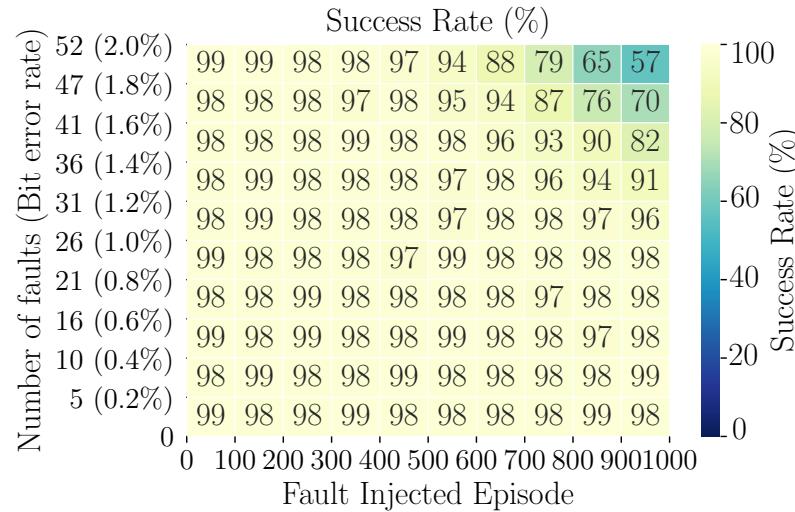


# Training – Transient Fault Impact



- Transient fault occurred in later episodes with high BER has higher impact

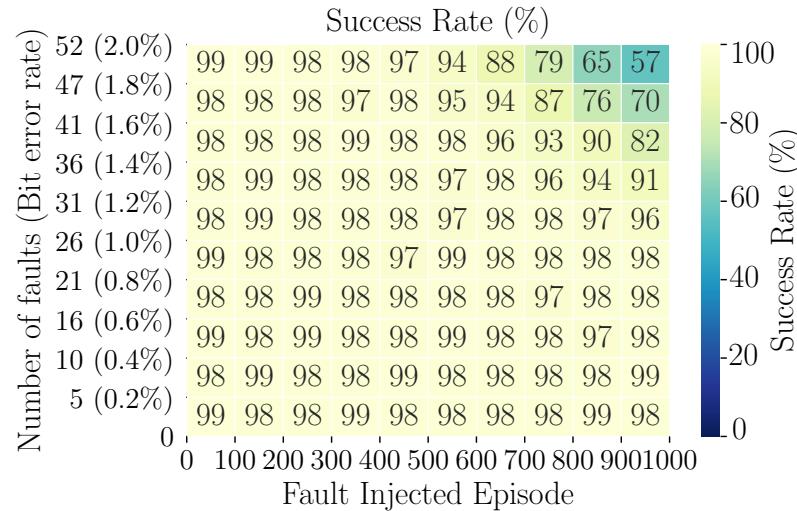
# Training – Transient Fault Impact



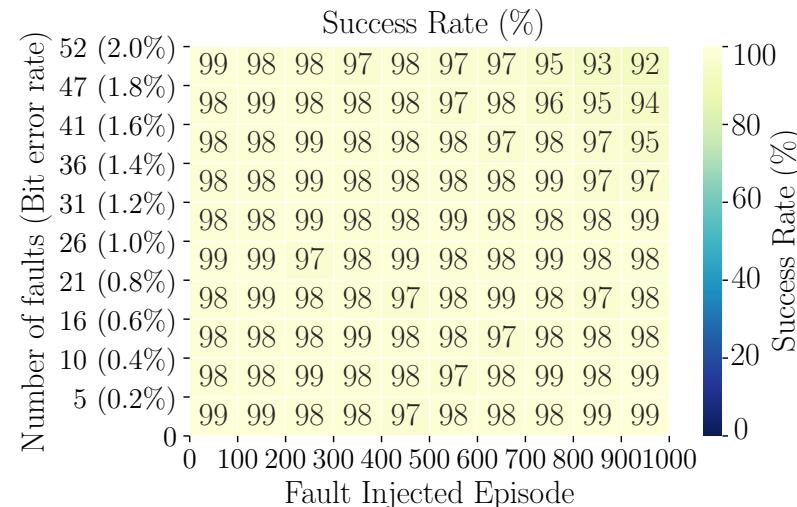
- **Transient fault occurred in later episodes with high BER has higher impact**
- **0 → 1 bit-flip has higher impact than 1 → 0 bit-flip**

# Training – Server and Agent Comparison

## Faults in server:



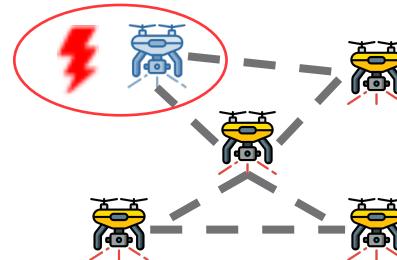
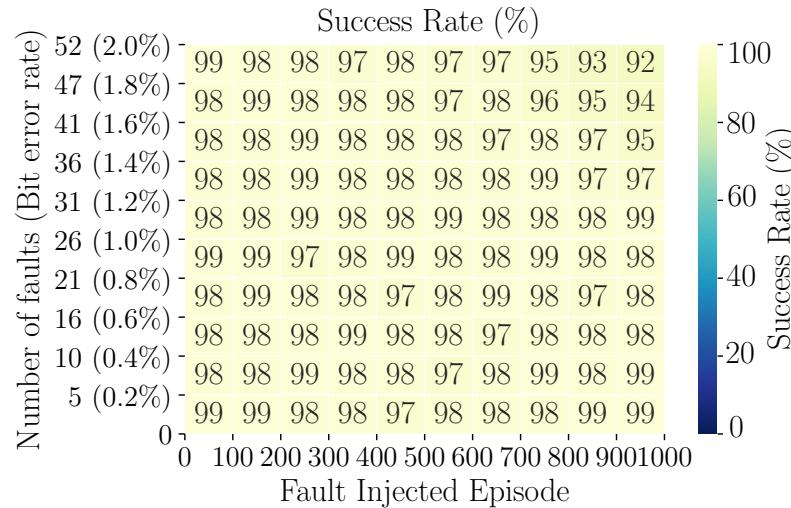
## Faults in agents:



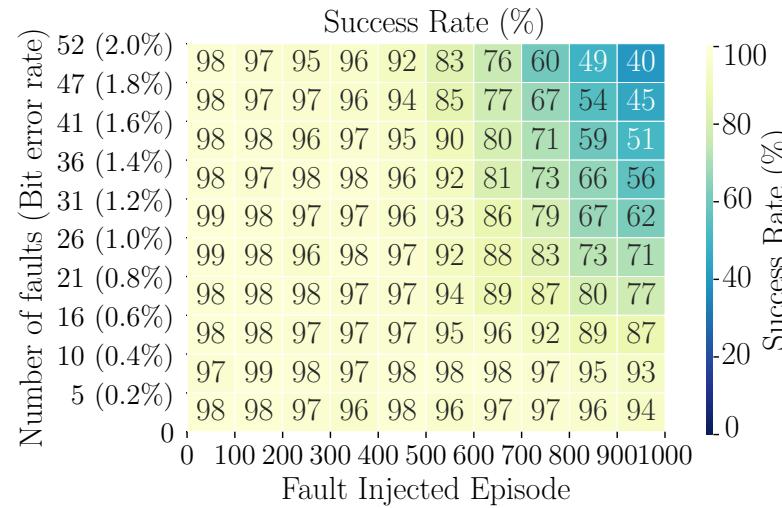
- **Server faults have higher impact than agent faults**
- **Apply fault mitigation method in server**

# Training – Single and Multi-Agent Comparison

Faults in agents of FRL system:

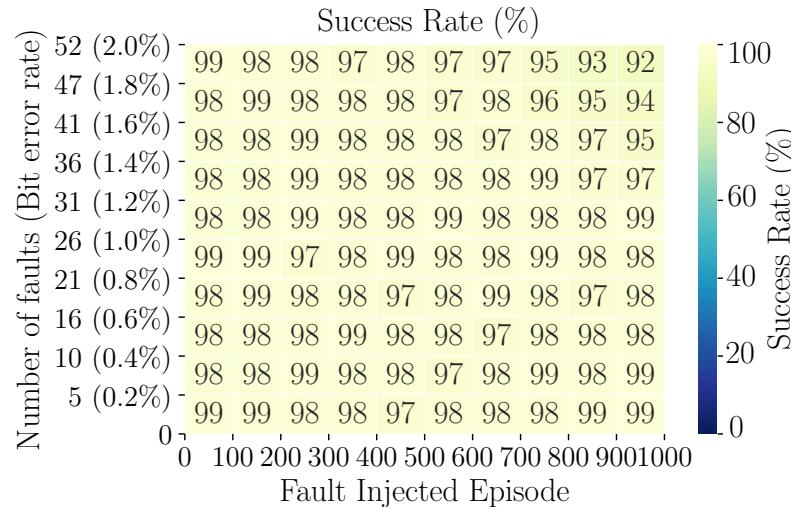


Faults in single-agent system:

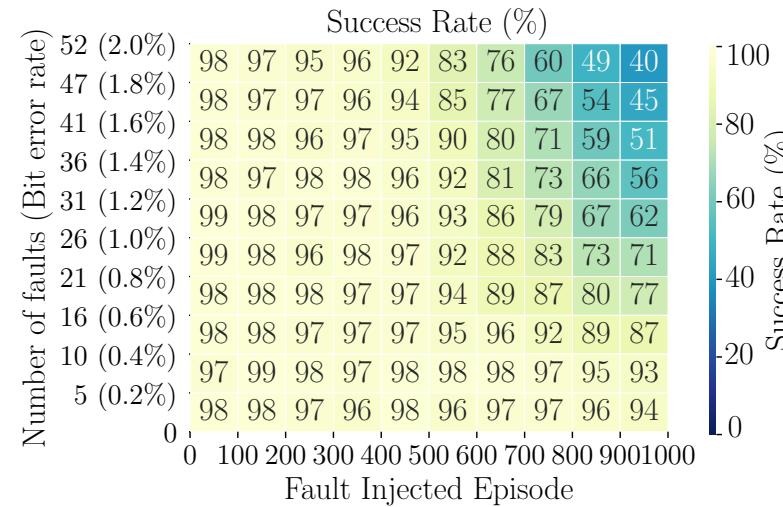


# Training – Single and Multi-Agent Comparison

Faults in agents of FRL system:



Faults in single-agent system:



- Multi-agent FRL system exhibits higher performance and resilience than single-agent system

# Training – Single and Multi-Agent Comparison

**TABLE I:** [GridWorld] Standard deviation ( $std$ ) of the consensus policy. Larger  $std$  indicates better differentiation between good and bad actions. Multi-agent system has higher  $std$  than single-agent system, indicating its higher performance and resilience.  $n$  means #agents.

	Single-agent	Multi-agent (n=4)	Multi-agent (n=8)	Multi-agent (n=12)
Std	0.255	0.405	0.472	0.504

(The larger of the value, the better of policy generalizability)

- Multi-agent FRL system exhibits higher performance and resilience than single-agent system
- Policy in multi-agent system is able to generalize better

# Training – Policy Convergence

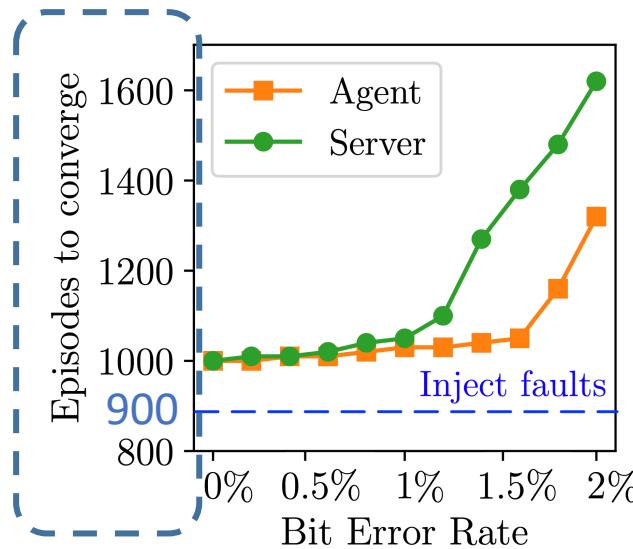
Question: Whether policy can finally converge after faults occurred?



- Faults are injected at 900<sup>th</sup> episode with different BER

# Training – Policy Convergence

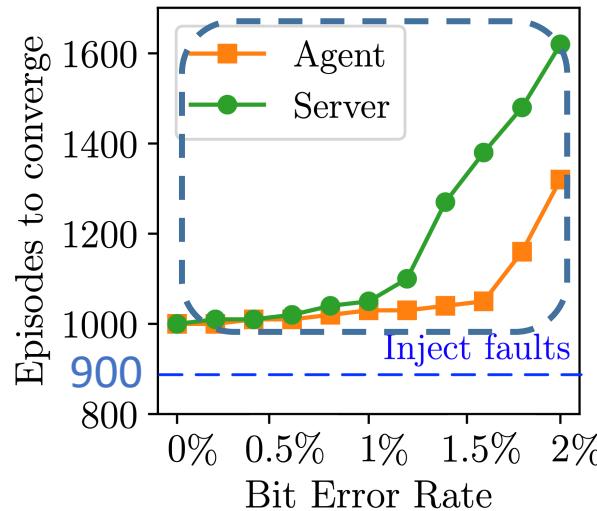
Question: Whether policy can finally converge after faults occurred?



- Faults are injected at 900<sup>th</sup> episode with different BER
- The episodes taken to converge after fault injected

# Training – Policy Convergence

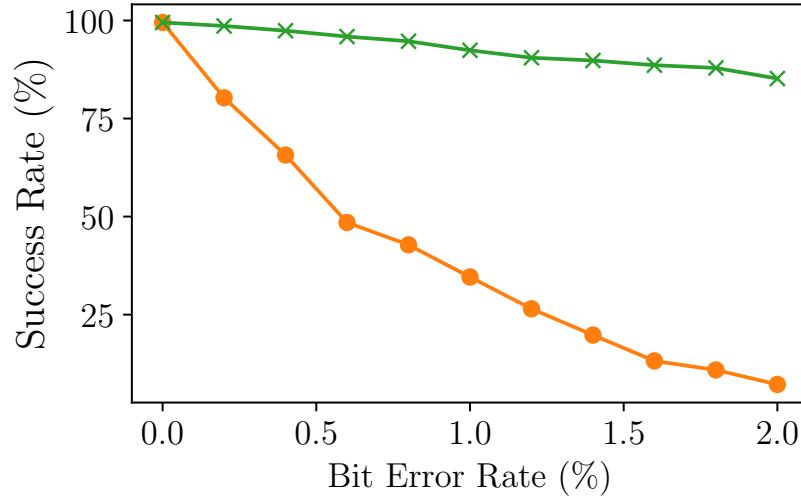
Question: Whether policy can finally converge after faults occurred?



- Faults are injected at 900<sup>th</sup> episode with different BER
- The episodes taken to converge after fault injected

- Transient faults will NOT affect policy convergence with longer fine-tuning training time

# Inference – Transient Fault Impact



Multi-Agent System:

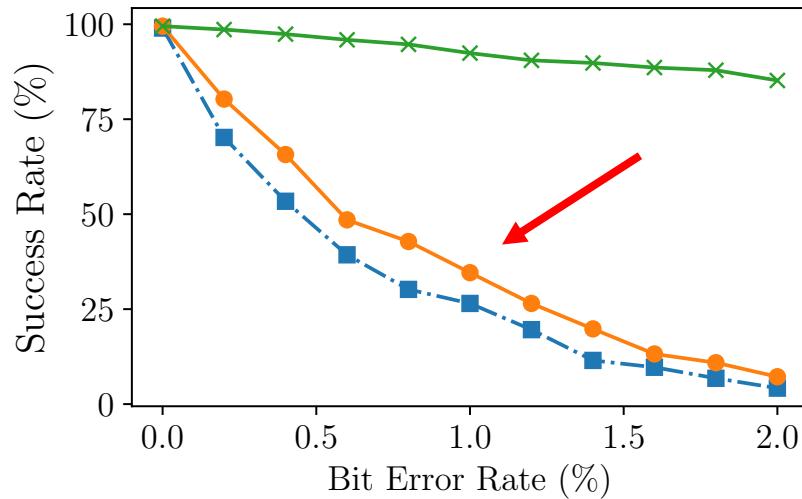
- Multi-Trans-M
- ×— Multi-Trans-1

Multi-Trans-M: faults affect all following action steps

Multi-Trans-1: faults affect one action step

- The sequential decision-making procedure of FRL system
- One action step fault does not necessarily result in task failure

# Inference – Single and Multi-Agent Comparison



Single-Agent System:

Single-Trans-M

Multi-Agent System:

Multi-Trans-M  
Multi-Trans-1

Multi-Trans-M:

faults affect all following action steps (FRL System)

Single-Trans-M:

faults affect all following action steps (Single-agent sys)

- Multi-agent FRL system is more resilient than single-agent system

# Drone Autonomous Navigation Problem

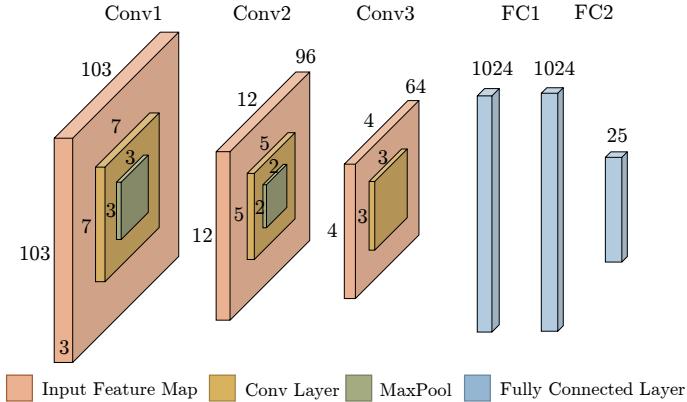
- Environments and Demos



(Powered by Unreal Engine and AirSim)



- Policy Architecture



- Evaluation Metric

- Drone safe flight distance  
(the longer, the better)

# Drone Autonomous Navigation Problem

- Environments and Demos



(Powered by Unreal Engine and AirSim)

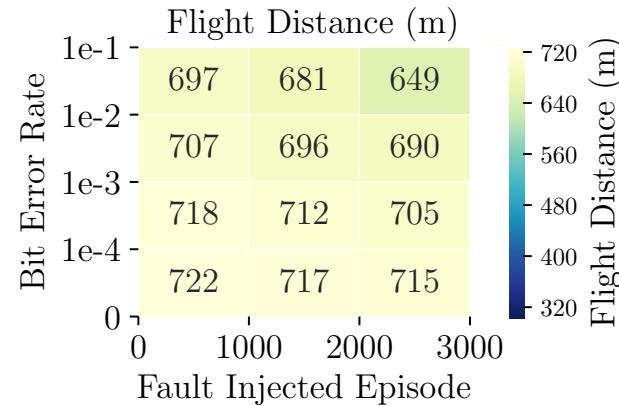


## Experiment Overview

- Training (on-device fine-tuning)
  - Transient fault impact
  - Single and multi-agent comparison
  - Different number of drones
  - Different communication intervals
- Inference
  - Different layer type
  - Different data type

# Training – Transient Fault Impact

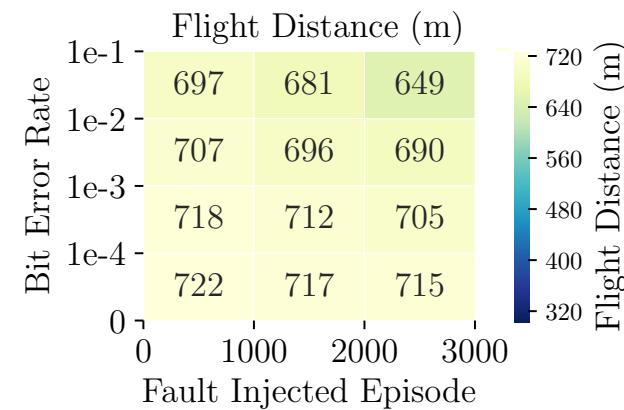
## FRL system: faults in agent



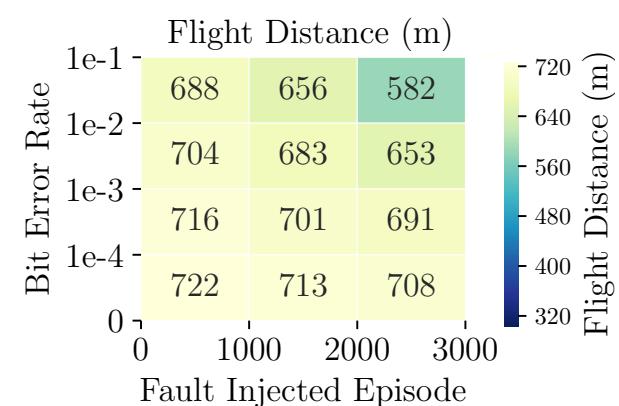
- **Faults occurred in later fine-tuning episodes with a higher BER impact the system more**

# Training – Single and Multi-Drone Comparison

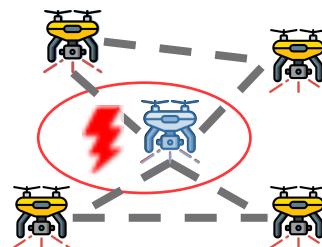
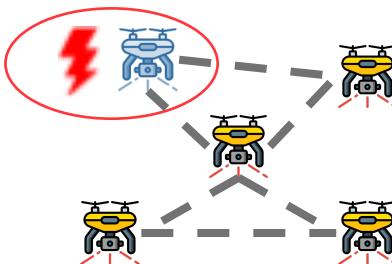
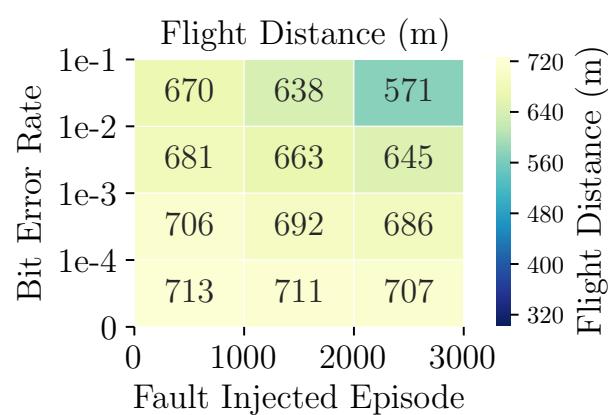
FRL system: faults in agent



FRL system: faults in server

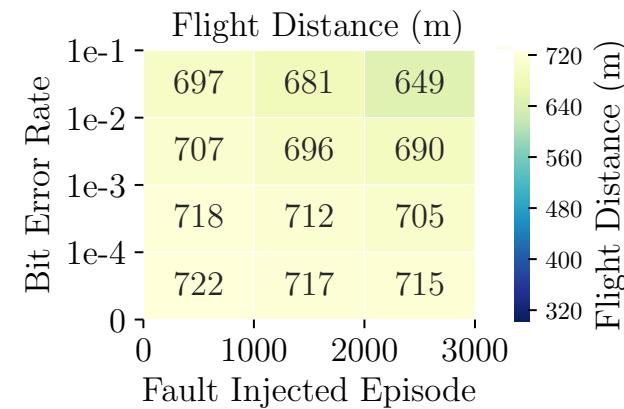


Single-agent system

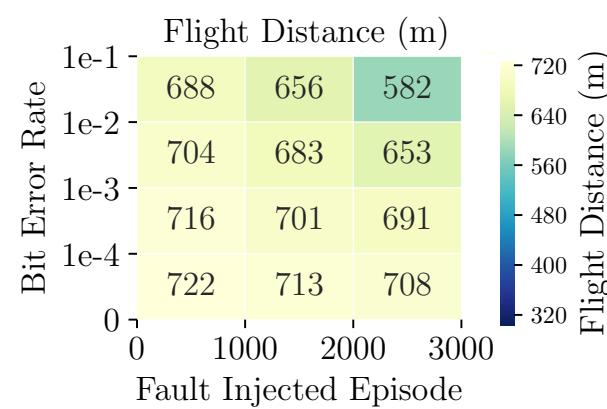


# Training – Single and Multi-Drone Comparison

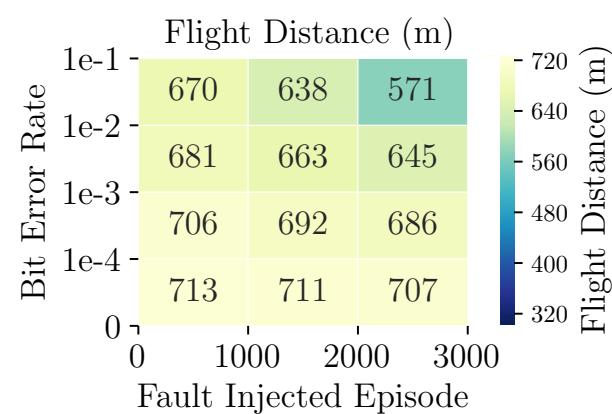
FRL system: faults in agent



FRL system: faults in server

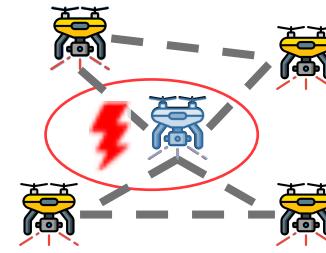
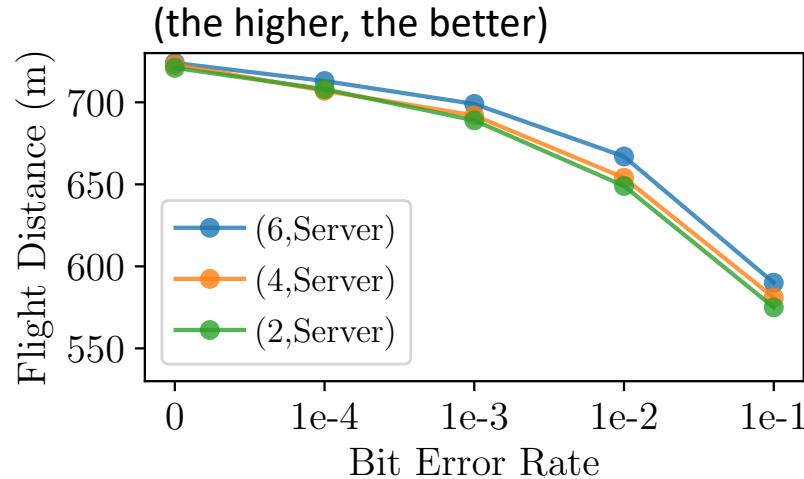


Single-agent system



- Multi-drone FRL system exhibit higher performance and resilience than single-drone system.

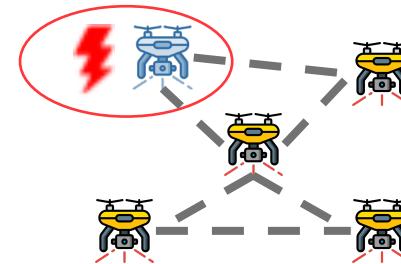
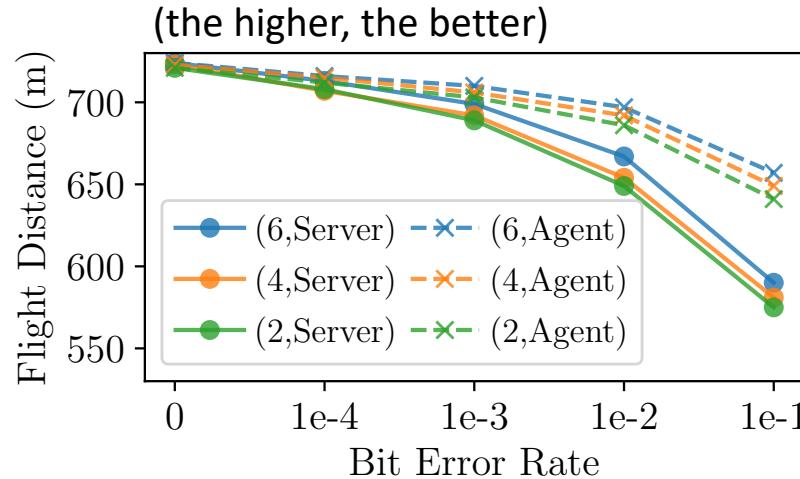
# Training – Number of Drones



(6/4/2, Server):  
(Total number of drones, faults in server)

- More drones helps improve FRL system resilience

# Training – Number of Drones

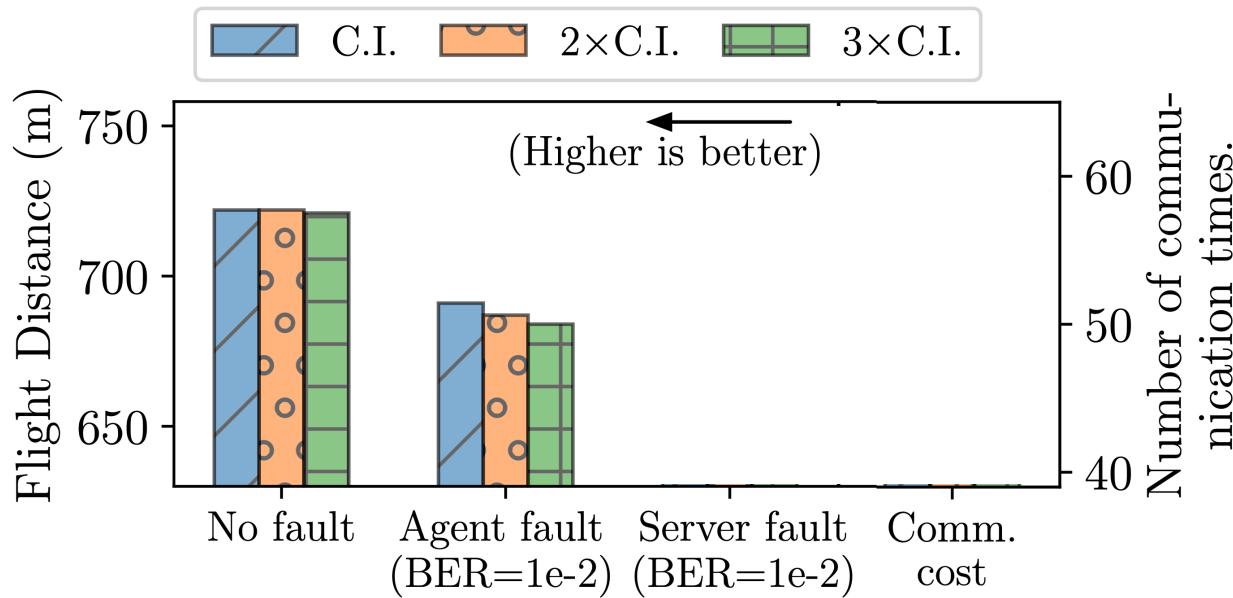


(6/4/2, Agent):  
(Total number of drones, faults in agent)

- More drones helps improve FRL system resilience

# Training – Communication Interval

(C.I.: Communication Interval)



After 2000<sup>th</sup> episode, drone perform more exploitation  
(policy is almost not updated)

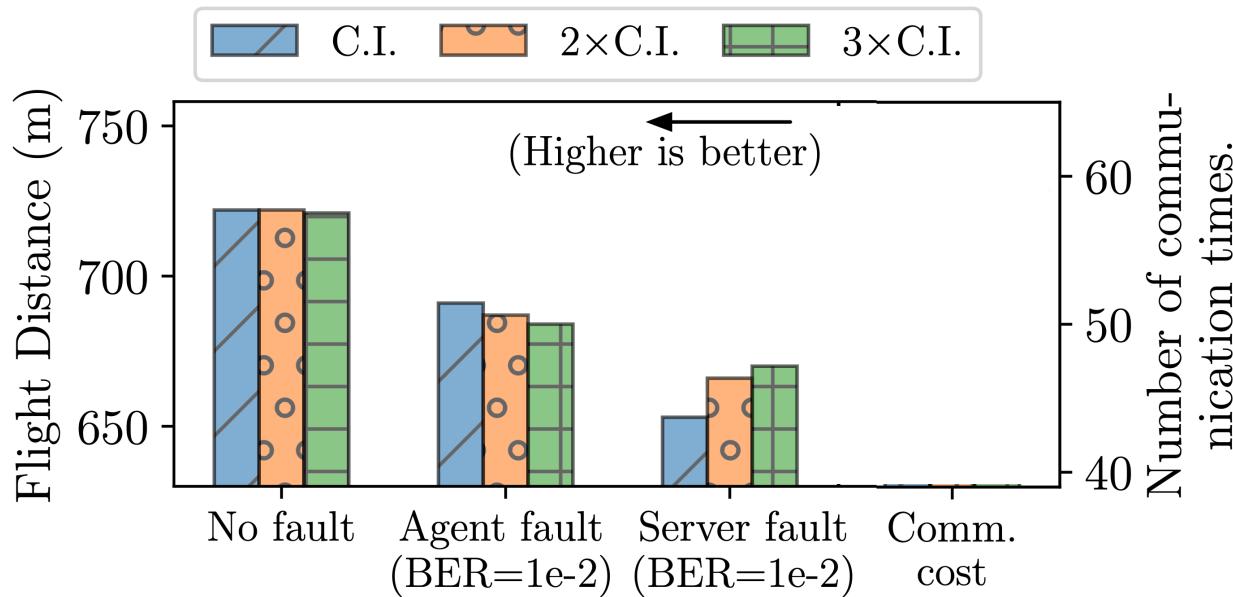


We increase the communication interval between agent and by 2x or 3x

- Longer communication interval makes FRL system more vulnerable to agent faults

# Training – Communication Interval

(C.I.: Communication Interval)



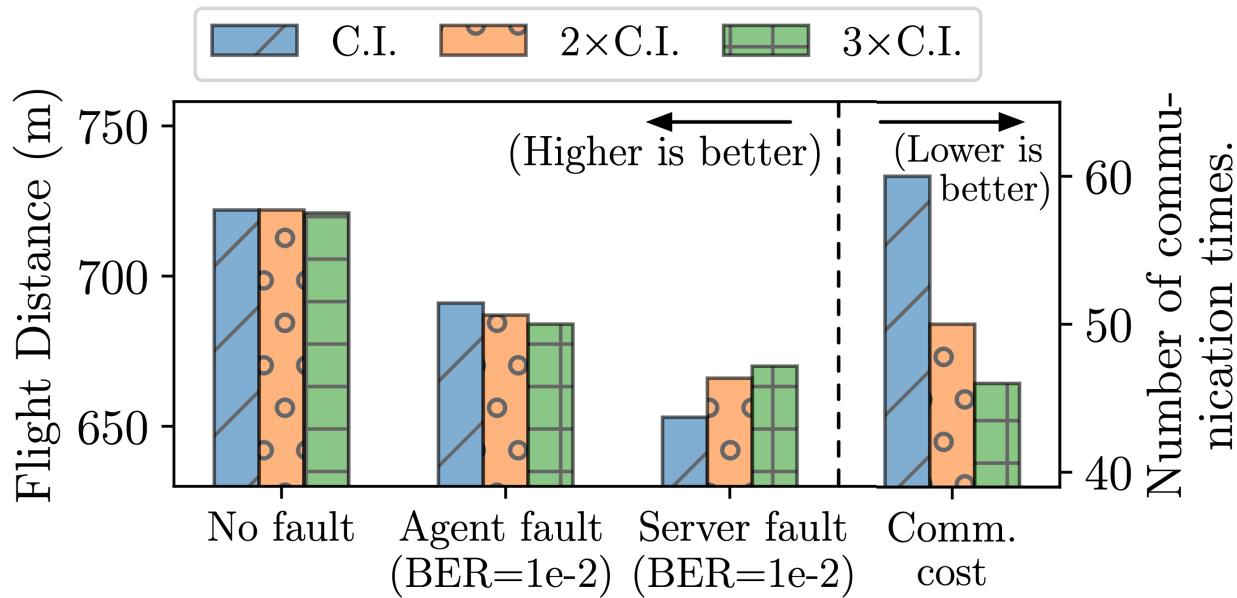
After 2000<sup>th</sup> episode, drone perform more exploitation  
(policy is almost not updated)

We increase the communication interval between agent and by 2x or 3x

- Longer communication interval alleviates impact of server fault

# Training – Communication Interval

(C.I.: Communication Interval)



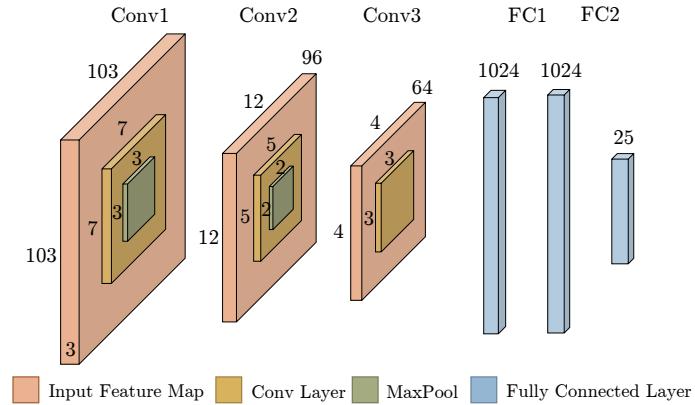
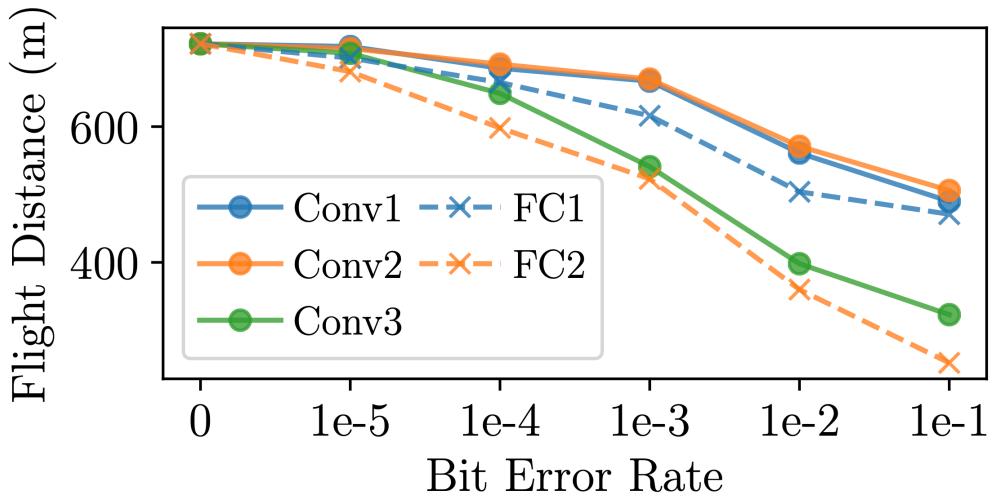
After 2000<sup>th</sup> episode, drone perform more exploitation  
(policy is almost not updated)



We increase the communication interval between agent and by 2x or 3x

- Longer communication interval reduce communication cost

# Inference – Layer Type



- Pooling and activation operations make layers more robust since bit-flips have higher probability of being masked and ceased propagation

# This Work

## Transient Fault Analysis for Federated Reinforcement Learning (FRL) -Based Navigation Systems



Transient fault injection for FRL-based navigation systems



Transient fault characterization for FRL-based navigation systems



Transient fault mitigation for FRL-based navigation systems

# Training – Server Checkpointing

- **Fault Indicator:** Cumulative reward drop exceeds  $x\%$  within  $y$  consecutive episodes ( $x$  and  $y$  are parameters)

## Fault in Agent

- **Detection:** agent  $i$  reward drop

## Fault in Server

- **Detection:** more than half of agents reward drop

# Training – Server Checkpointing

- **Fault Indicator:** Cumulative reward drop exceeds  $x\%$  within  $y$  consecutive episodes ( $x$  and  $y$  are parameters)
- **Fault Mitigation:** save checkpoint in server and update every 5 communication intervals

## Fault in Agent

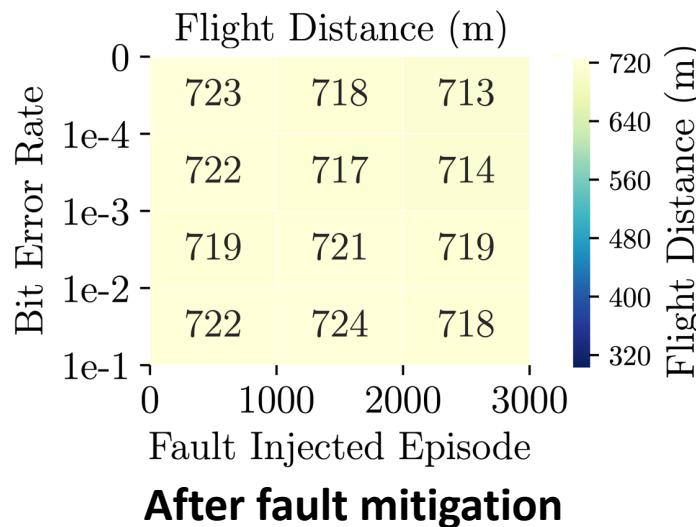
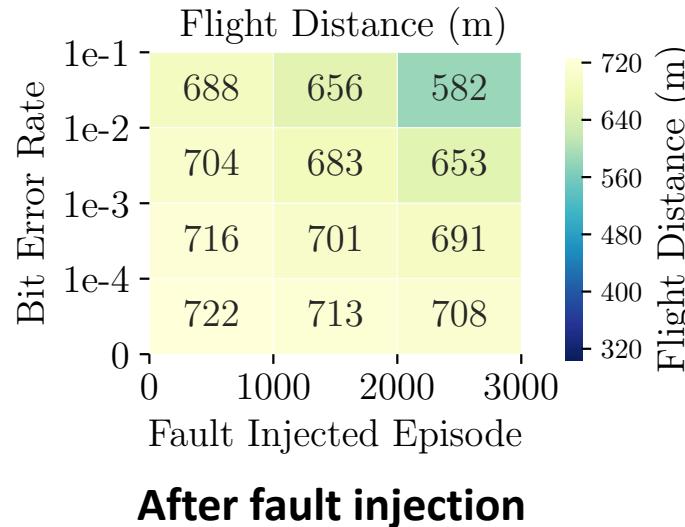
- Detection: agent  $i$  reward drop
- Recovery: copy server checkpoint to agent  $i$  memory

## Fault in Server

- Detection: more than half of agents reward drop
- Recovery: copy server checkpoint to server memory

# Training – Server Checkpointing

- **Evaluation**



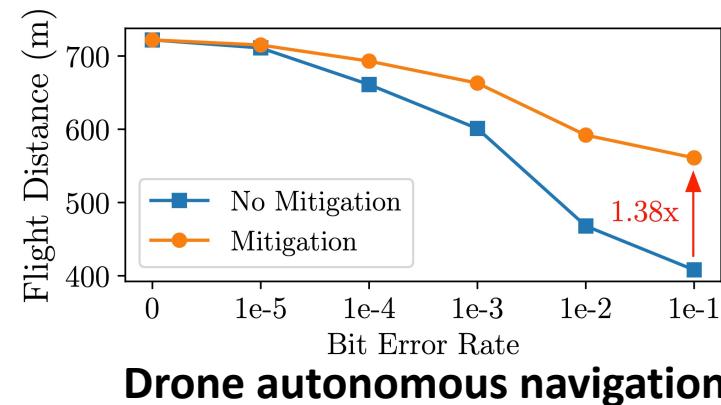
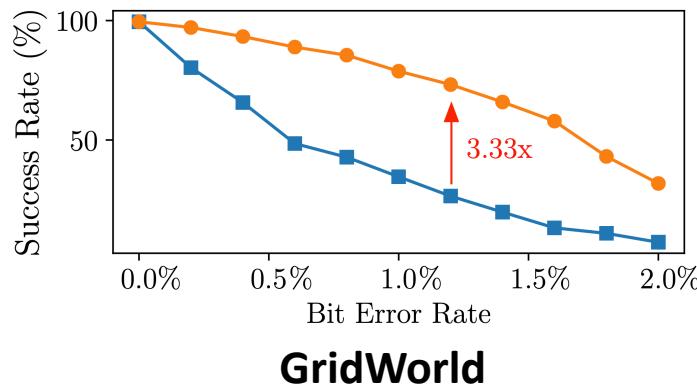
- **The impact of transient fault during training can be alleviated**

# Inference – Range-Based Anomaly Detection

- **Detection:** Statistically anomaly detection,  $(ai, bi) \rightarrow (1.1ai, 1.1bi)$
- **Recovery:** skip faulty operations

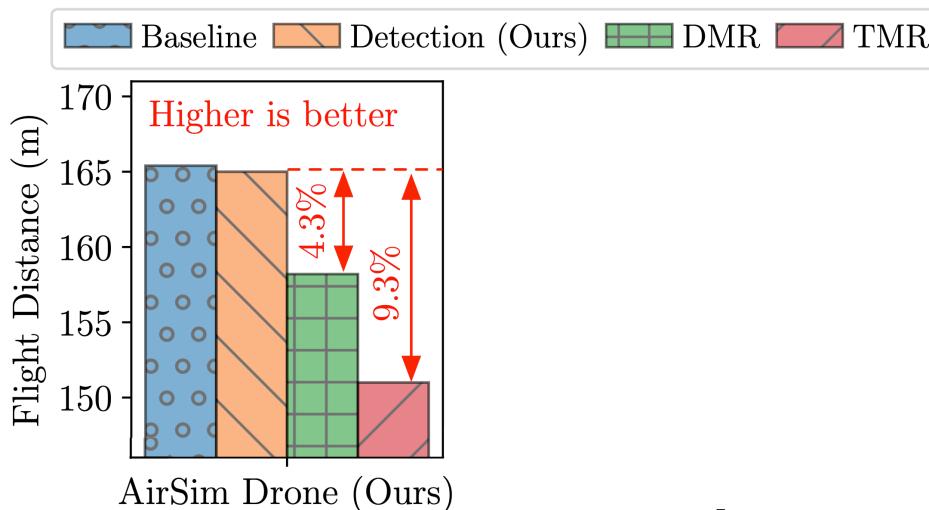
# Inference – Range-Based Anomaly Detection

- **Detection:** Statistically anomaly detection,  $(ai, bi) \rightarrow (1.1ai, 1.1bi)$
- **Recovery:** skip faulty operations
- **Evaluation:**



# Inference – Range-Based Anomaly Detection

- Compute overhead



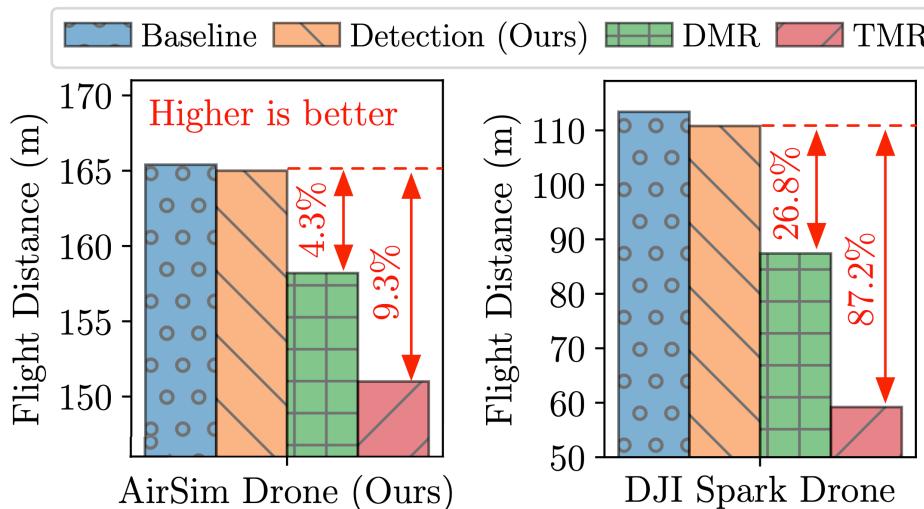
	AirSim Drone
Type	mini-UAV
Size (mm)	650
Weight (g)	1652
Battery Capacity (mAh)	6250



- Using a drone roofline-like performance model [Krishnan et al, CAL 2020]
- Negligible overhead compared to DMR, TMR

# Inference – Range-Based Anomaly Detection

- Compute overhead



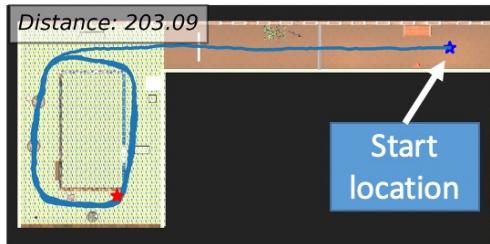
	AirSim Drone	DJI Spark
Type	mini-UAV	micro-UAV
Size (mm)	650	170
Weight (g)	1652	300
Battery Capacity (mAh)	6250	1480



- Using a drone roofline-like performance model [Krishnan et al, CAL 2020]
- Negligible overhead compared to DMR, TMR

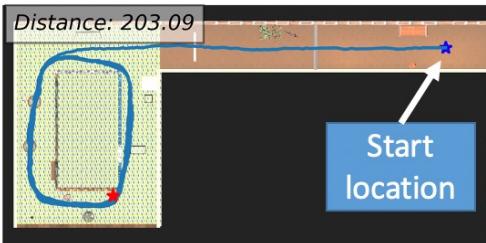
# Drone Flight Trajectory Demo

No Fault:

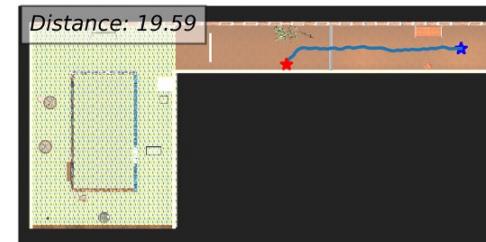
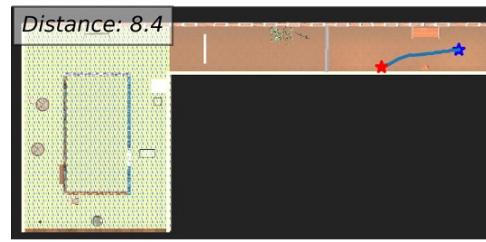
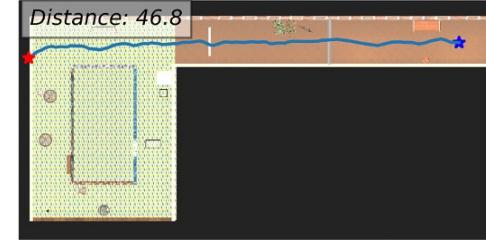
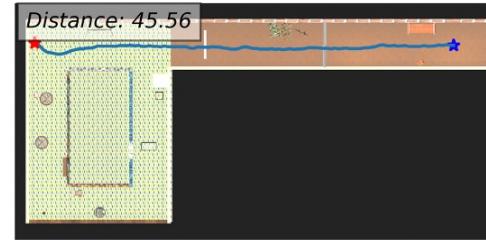


# Drone Flight Trajectory Demo

No Fault:

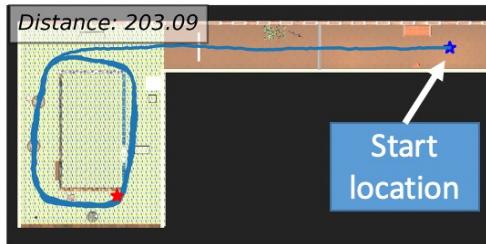


Fault injection:

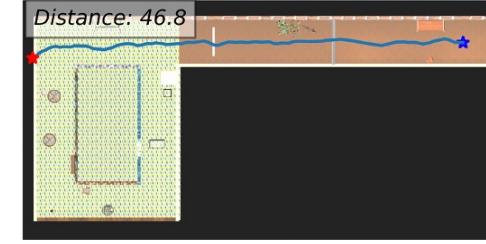
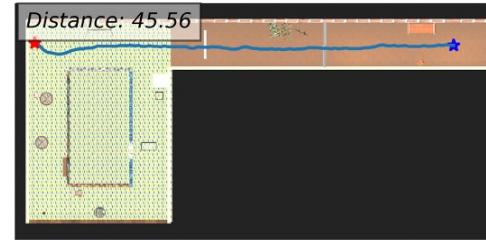


# Drone Flight Trajectory Demo

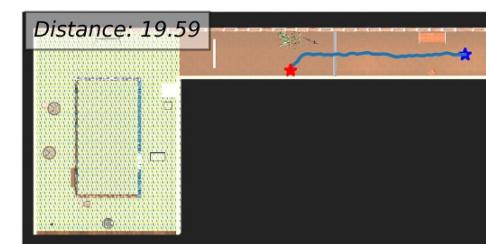
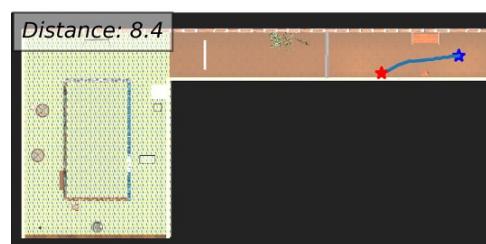
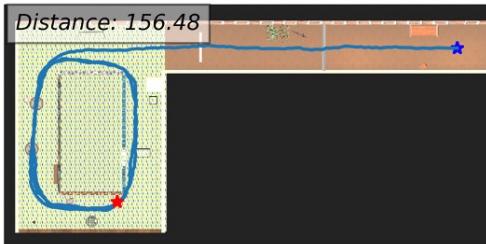
No Fault:



Fault injection:



Fault mitigation:



# In This Talk

## Transient Fault Analysis for Federated Reinforcement Learning (FRL) -Based Navigation Systems

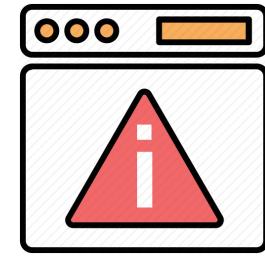


The **safety and reliability** of **swarm intelligence** **navigation systems** is important, but not well understood

A **fault injection** framework that emulates **hardware faults** and enables rapid fault analysis of FRL systems



Large-scale **fault injection study** in both training and inference stages of FRL systems against transient faults



Low-overhead **fault detection and recovery techniques** for both training and inference